

# Basic ODE fitting

November 12, 2018

## Contents

<b>1 Preliminaries</b>	<b>1</b>
<b>2 Basic fitting</b>	<b>1</b>
2.1 Exponential decay model . . . . .	1
2.2 Chemical reaction - multiple state fitting . . . . .	5
2.3 Fitting SIR model . . . . .	7

## 1 Preliminaries

Load packages:

```
library(fitode)
```

## 2 Basic fitting


### 2.1 Exponential decay model

Suppose we have a stock (compartment) that is losing a fixed *per capita* fraction per unit time, as in radioactive decay, financial depreciation, or a pure-death process.

```
set.seed(123)
true.m <- 0.5
true.A0 <- 200
true.sd <- 15

exp.data <- data.frame(
  x=rep(0:4, 3),
  y=rnorm(15, true.A0 * exp(-true.m * rep(0:4, 3)), sd=true.sd)
)
```

```
plot(exp.data)
```



The true dynamics can be modeled with the following equation:

$$\frac{dA}{dt} = -mA$$

We can translate this into a **fitode** model as follows:

```
exp.model <- new("model.ode",
  name = "SI",
  model = list(
    A ~ -m * A
  ),
  observation = list(
    y ~ dnorm(mean=A, sd=sd)
  ),
  initial = list(
    A ~ AO
  ),
  par=c("m", "AO", "sd")
)
```


Then, we can fit the model:

```
exp.fit <- fitode(
  exp.model,
  exp.data,
  start=c(m=0.5, A0=200, sd=15),
  tcol="x"
)

## Fitting ode ...
## Computing vcov on the original scale ...
```

To diagnose the fit, we can use the `plot` method for `fitode` objects. *bmb: we should be careful here: in base R the convention is that `plot` shows diagnostics (although of course that's different in `deSolve`, which plots trajectories). Do we want to set up different methods for plotting predictions vs trajectories, or should there be an option to `plot`?* By default, this will plot the estimated trajectory along with 95% confidence intervals, estimated via the delta method.

```
plot(exp.fit, level=0.95)
curve(true.A0 * exp(-true.m*x), add=TRUE, lty=1, col="red")
legend(
  x="topright",
  legend=c("estimated", "true"),
  col=c("black", "red"),
  lty=1
)
```



To obtain confidence intervals on the parameters, we can use `confint` method. There are three available methods for obtaining the confidence intervals: `delta`, `profile` and `wmvrnorm`. Due to computation speed, default option is `delta`. We will get into the details later. For now,

```
confint(exp.fit)

##           estimate      2.5 %      97.5 %
## m      0.5303757    0.4757616    0.5912593
## A0 212.1596895 200.6325622 224.3490954
## sd  11.0526632   7.7278640  15.8079080
```

Since this particular differential equation has a simple closed-form solution ( $A(t) = A(0) \exp(-mt)$ ), we can use the `glm` function to fit the model as well (we need to use `glm` with the `gaussian` family and a log link, rather than `lm`, in order to fit an exponential curve with constant residual variance).


```
glm.fit <- glm(y~x,
  family=gaussian(link="log"),
  data=exp.data,
  start = c(intercept=log(200), x=-0.5))
```

Comparing the results:

```
glm.pred <- predict(glm.fit, data.frame(x=0:4), se.fit=TRUE, type="response")

glm.data <- data.frame(
  x=0:4,
  estimate=glm.pred$fit,
  lwr=glm.pred$fit-1.96 * glm.pred$se.fit,
  upr=glm.pred$fit+1.96 * glm.pred$se.fit
)


plot(exp.fit, level=0.95)
lines(glm.data$x, glm.data$estimate, col=2)
lines(glm.data$x, glm.data$lwr, col=2)
lines(glm.data$x, glm.data$upr, col=2)
legend(
  x="topright",
  legend=c("estimated (fitode)", "estimated (glm)"),
  col=c("black", "red"),
  lty=1
)
```



The estimated trajectories and their confidence intervals are essentially indistinguishable.

## 2.2 Chemical reaction - multiple state fitting

Now, consider the following chemical reaction:



Then, we can write the governing differential equations as follows:

$$\begin{aligned}\frac{dA}{dt} &= -kA \\ \frac{dB}{dt} &= 3kB\end{aligned}$$

Suppose we have observed values for both quantities (we assume their standard deviation is identical):

Taking a quick look at the simulation results:

```
head(reaction_data)
##      times      y1      y2
## 1      1 317.8691 106.8864
## 2      2 276.4298 191.1854
## 3      3 225.9531 262.5231
## 4      4 229.2591 330.2038
## 5      5 196.3682 392.9070
## 6      6 171.2810 447.1751
```

Here,  $y_1$  measures quantity  $A$  and  $y_2$  measures quantity  $B$ . *bmb: is there a reason not to use names that match the ODE definition?* Then, we can define the model:


```
reaction_model <- new("model.ode",
  name = "reaction",
  model = list(
    A ~ - k * A,
    B ~ 3 * k * A
  ),
  observation = list(
    y1 ~ dnorm(mean=A, sd=sd),
    y2 ~ dnorm(mean=B, sd=sd)
  ),
  initial = list(
    A~A0,
    B~B0
  ),
  par=c("k", "A0", "B0", "sd")
)
```

We can fit this using arbitrary starting conditions.

```
reaction_fit <- fitode(
  reaction_model,
  reaction_data,
  start=c(k=0.1, A0=300, B0=10, sd=10)
)

## Fitting ode ...
## Computing vcov on the original scale ...

plot(reaction_fit, level=0.95)
```




Confidence intervals...

```
confint(reaction_fit)
```

	estimate	2.5 %	97.5 %
## k	0.100314	0.09714102	0.1035907
## A0	299.581692	295.20809184	304.0200894
## B0	100.734330	90.39960933	112.2505441
## sd	9.151578	7.35070035	11.3936588

## 2.3 Fitting SIR model

```
harbin <- fitsir::harbin
harbin2 <- rbind(data.frame(week=1, Deaths=NA), harbin)
plot(harbin2)
```



We need to add an NA observation to make this work...

```
SI_model_c <- new("model.ode",
  name = "SI",
  model = list(
    S ~ - beta*S*I/N,
    I ~ beta*S*I/N - gamma*I,
    cDeath ~ gamma*I
  ),
  observation = list(
    Deaths ~ dnbinom(mu=cDeath, size=size)
  ),
  initial = list(
    S ~ N * (1 - i0),
    I ~ N * i0,
    cDeath ~ 0
  ),
  diffnames="cDeath",
  par=c("beta", "gamma", "N", "i0", "size")
)

start <- c(beta=2, gamma=1, N=20000, i0=1e-5, size=10)

sirfit <- fitode(
  SI_model_c,
  harbin2,
```




```

start=start,
link = list(
  beta="log",
  gamma="log",
  N="log",
  i0="logit",
  size="log"
),
tcol="week"
)

## Fitting ode ...
## Computing vcov on the original scale ...

plot(sirfit, level=0.95)

```



Confidence intervals on various epidemiological quantities:

```

confint(sirfit,
  parm=list(
    R0~beta/gamma,
    r~beta-gamma
  ),
  method="wmvrnorm")

##      estimate      2.5 %    97.5 %
## R0  1.705258  1.4733837  2.132810

```


```
## r 0.772026 0.6942409 0.837602
```

Alternate parameterization:

```
SI_model_c2 <- Transform(  
  SI_model_c,  
  list(  
    beta~(R0_1+1)*gamma  
  ),  
  par=c("R0_1", "gamma", "N", "i0", "size")  
)  
  
cc <- coef(sirfit)  
  
start2 <- c(R0_1=unname(cc[1]/cc[2]-1), cc[-1])  
  
sirfit2 <- fitode(  
  SI_model_c2,  
  harbin2,  
  start=start2,  
  link = list(  
    R0_1="log",  
    gamma="log",  
    N="log",  
    i0="logit",  
    size="log"  
  ),  
  tcol="week"  
)  
  
## Fitting ode ...  
## Computing vcov on the original scale ...
```

Compare fits:

```
plot(sirfit)  
plot(sirfit2, add=TRUE, col.traj="red", col.conf="red")
```



We get identical fits. The advantage of this parameterization is that we can obtain profile confidence intervals on  $R_0$  (it's a little slow...):

```
set.seed(101)
confint(sirfit2, "R0_1", method="profile") + 1

##      estimate    2.5 %   97.5 %
## R0_1 1.705259 1.000006 2.591917

confint(sirfit2, "R0_1", method="wmvnorm") + 1

##      estimate    2.5 %   97.5 %
## R0_1 1.705259 1.238815 2.373596

confint(sirfit, parm=list(R0~beta/gamma))

##      estimate    2.5 %   97.5 %
## R0 1.705258 0.9361381 2.474379

confint(sirfit, parm=list(R0~beta/gamma), method="wmvnorm")

##      estimate    2.5 %   97.5 %
## R0 1.705258 1.47908 2.147964
```

I'm not sure why performing `wmvnorm` on the  $\beta/\gamma$  scale gives narrower confidence intervals. Maybe it's because of the parameterization?

```
set.seed(101)
plot(sirfit, method="wmvnorm")
plot(sirfit2, add=TRUE, col.traj="red", col.conf="red", method="wmvnorm")
```

